

Limitations / Known Issues

- The application relies on Large Language Models (Gemini) for research analysis, summaries, and insights. Generated responses may occasionally contain inaccuracies or omissions.
- PDF text extraction quality depends on the document format. Scanned or image-based PDFs without selectable text may produce incomplete results.
- The citation extraction agent works best with well-structured academic papers. Papers with unusual formatting may not have all references extracted correctly.
- Review scores are generated using a combination of LLM evaluation and rule-based validation. The score is intended as a quality indicator rather than an absolute measure of research quality.
- Public PDF URLs (such as arXiv or direct PDF links) are supported. Some publisher websites (IEEE Xplore, ScienceDirect, Springer, ACM, etc.) may require authentication or institutional access, in which case users should upload the PDF directly.
- Large research papers may require additional processing time depending on the paper length and LLM response latency.
- The system currently processes one research paper at a time and does not support batch analysis.
- Metadata such as publication year or venue may not always be available if it is missing from the uploaded paper.
- The application requires a valid Gemini API key and an active internet connection for AI-based analysis.